

# LINFO1104

## Concepts, paradigms, and semantics of programming languages

### Lecture 10-11

Peter Van Roy

ICTEAM Institute  
Université catholique de Louvain

[peter.vanroy@uclouvain.be](mailto:peter.vanroy@uclouvain.be)



1

## Overview of lectures 10-11

- Limitation of deterministic dataflow
  - Some applications cannot be written in deterministic dataflow! We explain why not.
- Message-passing concurrency (multi-agent actor programming)
  - We overcome the limitation by adding a new concept, **ports**, to deterministic dataflow
  - We define **port objects** and **active objects**
  - We show how to write message protocols and large multi-agent actor programs



2

# Limitation of deterministic dataflow



3

## Limitation of deterministic dataflow



- In lectures 8-9 we saw deterministic dataflow, which makes concurrent programming very easy
  - It allows “Concurrency for Dummies”: threads can be added to the program at will without changing the result
- But unfortunately it cannot be used all the time!
  - It has a strong limitation: it cannot be used to write programs when the nondeterminism must be visible
  - But why must nondeterminism sometimes be visible? Let’s see an example: a client/server application.

4


## Client/server application (1)



- A client/server application consists of a set of clients all communicating to one server
  - The clients and the server are concurrent agents
  - Each client sends messages to the server and receives replies
- Client/server applications are ubiquitous on the Internet
  - For example, all Web stores are client/servers: the users are clients and the store is the server
  - When shopping at Amazon, your Web browser sends messages and receives replies from the Amazon server
- Client/server cannot be written in deterministic dataflow!
  - Why not? Let's try and see what goes wrong! Try it yourself!

5

## Client/server application (2)



- Each client has a link to the server and can send messages to the server at any time
- The server receives each message, does a local computation, and then replies immediately

6

## Client/server: first attempt



- Let's try to write a client/server in deterministic dataflow
  - Assume that there are two clients, each with an output stream, and the server receives both
- Here is the server code:

```
proc {Server S1 S2}
  case S1|S2 of (M1|T1)|S2 then
    (handle M1) {Server T1 S2}
  [] S1|(M2|T2) then
    (handle M2) {Server S1 T2}
  end
end
```

This doesn't work!  
Explain why not.

7

## Client/server: second attempt



- The first attempt does not work if Client 2 sends a message and Client 1 sends nothing
- We can try doing it the other way around:

```
proc {Server S1 S2}
  case S1|S2 of S1|(M2|T2) then
    (handle M1) {Server S1 T2}
  [] (M1|T1)|S2 then
    (handle M2) {Server T1 S2}
  end
end
```

- This doesn't work if Client 1 sends a message and Client 2 sends nothing!

8

## Client/server: third attempt



- Maybe the server has to receive from both clients:

```
proc {Server S1 S2}
  case S1|S2 of (M1|T1)|(M2|T2) then
    (handle M1)
    (handle M2) {Server T1 T2}
  end
end
```

- This does not work either! (Why not?)

9

## What is the problem?



- The **case** statement waits on **a single pattern**
  - This is because of determinism: with the same input, the **case** statement must give the same result
- But the server must wait on **two patterns**
  - Either M1 from Client 1 or M2 from Client 2
  - Either pattern is possible, it depends on when each client sends the message and on how long the message takes to reach the server
    - The decision is made **outside the program**
  - This means exactly that execution is nondeterministic!

10

## Understanding nondeterminism



- Nondeterminism means that a choice is made **outside of the program's control**
  - This is exactly what is happening here: the choice is the arrival order of the client messages, which depends on the human clients and on the message travel time
- The nondeterminism is inherently part of the client/server execution, it cannot be avoided
  - The nondeterminism is a consequence of the initial requirement: "The server receives each message, does a local computation, and then replies immediately"
  - This means that the reply cannot be delayed while the server waits for another message

11

## Overcoming the limitation



12

## Overcoming the limitation



- Deterministic dataflow cannot express an application that requires nondeterminism
- To do this, we need to extend the kernel language with a new concept
- The new concept must be able to wait on several events nondeterministically
  - The new language is no longer deterministic!
- We will show two possible solutions

13

## Solution 1: WaitTwo



- We introduce the function:  
    {WaitTwo X Y}  
    with the following semantics:
  - {WaitTwo X Y} can return 1 if X is bound
  - {WaitTwo X Y} can return 2 if Y is bound
  - If either X or Y is bound, {WaitTwo X Y} will return
- If both X and Y are unbound, it just waits
- If both X and Y are bound, it can return either 1 or 2, both are possible (nondeterminism!)

14

## Client/Server with WaitTwo



- Here is the client/server with WaitTwo:

```
proc {Server S1 S2}
  C={WaitTwo S1 S2}
in
  case C|S1|S2 of 1|(M1|T1)|S2 then
    (handle M1) {Server T1 S2}
  [] 2|S1|(M2|T2) then
    (handle M2) {Server S1 T2}
  end
end
```

- If Client 1 sends a message, C=1 and it is handled
- If Client 2 sends a message, C=2 and it is handled
- What happens if both Client 1 and Client 2 send messages?

15

## WaitTwo is not scalable



- What happens if we have millions of clients?
  - WaitTwo solves the problem for two clients
  - How can we wait on millions of clients?
- One possibility is to “merge” all client streams into a single stream:

```
fun {Merge S1 S2}
  C={WaitTwo S1 S2}
in
  case C|S1|S2 of 1|(M1|T1)|S2 then M1|{Merge T1 S2}
  [] 2|S1|(M2|T2) then M2|{Merge S1 T2}
  end
end
```

- With Merge we build a huge tree of stream mergers. It must expand and contract if new clients arrive or old clients leave. Not very nice!

16



## Solution 2: Ports



- A better solution is to add ports (named streams)
- Ports have two operations:  

```
P={NewPort S} % Create port P with stream S
{Send P X}      % Add X to end of port P's stream
```
- How does this solve our problem?
  - With a million clients  $C_1$  to  $C_{1000000}$ :  
Each client  $C_i$  does  $\{\text{Send } M_i \text{ P}\}$  for each message it sends
  - The server reads the stream S, which contains all messages from all clients in some nondeterministic order

17

## Port example operations



- We create a port and do sends:  

```
P={NewPort S}
{Browse S} % Displays _
{Send P a} % Displays a|_
{Send P b} % Displays a|b|_
```
- What happens if we do:  

```
thread {Send P c} end
thread {Send P d} end
```
- What are the possible results of these two sends for all choices of the scheduler?

18

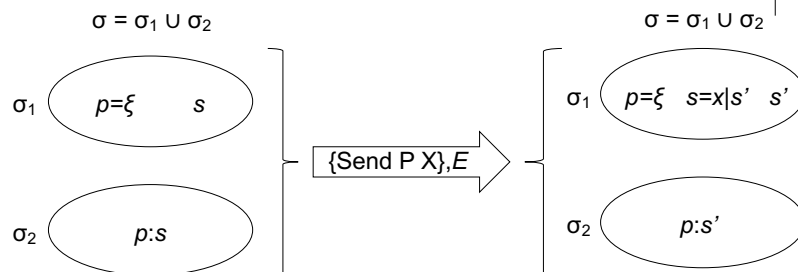
## Port semantics (1)



- Assume single-assignment store  $\sigma_1$  with variables
- Assume a port store  $\sigma_3$  that contains pairs of variables
- $P = \{\text{NewPort } S\}, \{P \rightarrow p, S \rightarrow s\}$  <sup>environment</sup>
  - Assume unbound variables  $p, s \in \sigma_1$
  - Create fresh name  $\xi$ , bind  $p = \xi$ , add pair  $p:s$  to  $\sigma_2$
- $\{\text{Send } P \ X\}, \{P \rightarrow p, X \rightarrow x\}$ 
  - Assume  $p = \xi$ , unbound variable  $s \in \sigma_1, p:s \in \sigma_2$
  - Create fresh unbound variable  $s'$ , bind  $s = x|s'$ , update pair to  $p:s'$

19

## Port semantics (2)



- $\{\text{Send } X \ P\}$  adds  $x$  to the end of the port's stream and updates the new end of stream
  - The send operation is **atomic**, which means the scheduler is guaranteed never to stop in the middle, so it happens as if it is **one indivisible step**
- We assume that environment  $E = \{P \rightarrow p, X \rightarrow x\}$

20

## Client/server with ports



- Assume port  $P = \{\text{NewPort } S\}$
- Client code: (any number of clients!)
  - Do  $\{\text{Send } P \ M\}$  to send message to server
- Server code:

```
proc {Server S}
  case S of M|T then
    (handle M) {Server T}
  end
end
```

21

## Message-passing concurrency



22

## Message-passing concurrency



- Message-passing concurrency is a new paradigm for concurrent programming
  - It consists of deterministic dataflow with ports
  - It is also called **multi-agent actor programming**
- We will show how to write concurrent programs in this new paradigm
  - We will define **port objects** and **active objects**
  - We show how to do **message protocols**
  - We show how to write **large multi-agent programs**

23

## Stateless port objects (stateless agents)



24

## Stateless port objects



- A **stateless port object** is a combination of a port, a thread, and a recursive list function
  - We also call it a **stateless agent**
- Each agent is defined in terms of how it replies to messages
- Each agent has its own thread, so there are no problems with concurrency
- Agents are a very useful concept!

25

## A math agent



- First the computation procedure:

```
proc {Math M}
  case M
  of add(N M A) then A=N+M
  [] mul(N M A) then A=N*M
  ...
end
end
```

26

## Making it a port object



- We add a port, a thread, and a recursive procedure:

```
MP={NewPort S}  
proc {MathProcess Ms}  
  case Ms of M|Mr then  
    {Math M} {MathProcess Mr}  
  end  
end  
thread {MathProcess S} end
```

27

## Using ForAll



- We replace MathProcess by ForAll:

```
proc {ForAll Xs P}  
  case Xs of nil then skip  
  [] X|Xr then {P X} {ForAll Xr P}  
  end  
end
```

- Using ForAll, we get:

```
proc {MathProcess Ms} {ForAll Ms Math} end
```

28

## Defining new port objects (1)



- A generic way to build stateless port objects:

```
fun {NewAgent0 Process}
  Port Stream
in
  Port={NewPort Stream}
  thread {ForAll Stream Process} end
Port
end
```

29

## Defining new port objects (2)



- A generic way to build stateless port objects:

```
fun {NewAgent0 Process}
  Port Stream
in
  Port={NewPort Stream}
  thread for M in Stream do {Process M} end end
Port
end
```

Using syntax  
of **for** loops


30

# Stateful port objects (stateful agents)



31

## Stateful port objects (Section 5.2)



- A stateful port object, also called stateful agent, has an internal memory  $s_i$  called its **state**
- The state is updated with each message received, which gives a **state transition function**:  
 $F: \text{State} \times \text{Msg} \mapsto \text{State}$

32



## Creating stateful port objects



- We define a generic function for stateful port objects:

```
fun {NewPortObject Init F}  
  proc {Loop S State}  
    case S of M|T then {Loop T {F State M}} end  
  end  
  P  
in  
  thread S in P={NewPort S} {Loop S Init} end  
  P  
end
```

33

## Structure of Loop



- Does the Loop function ring a bell?

```
proc {Loop S State}  
  case S of M|T then {Loop T {F State M}} end  
end
```

- Loop starts from an initial state
- Loop successively applies F to the previous state and a message
- The function F is a binary operation
- ...

34

## Structure of Loop



- Does the Loop function ring a bell?

```
proc {Loop S State}
  case S of M|T then {Loop T {F State M}} end
end
```

- Loop starts from an initial state
- Loop successively applies F to the previous state and a message
- The function F is a binary operation
- Of course! It is a Fold operation!

35

## FoldL operation



- FoldL is an important higher-order operation:

```
fun {FoldL S F U}
  case S
  of nil then U
  [] H|T then {FoldL T F {F U H}}
  end
end
```

36

## Fold is the heart of the agent



- We replace:  
`thread S in P={NewPort S} {Loop S Init} end`
- by:  
`thread S in P={NewPort S} {FoldL S F Init} end`
- Oops! There is a small bug...

37

## Updated NewPortObject



- We define a generic function for stateful port objects:  

```
fun {NewPortObject Init F}  
  P Out  
in  
  thread S in P={NewPort S} Out={FoldL S F Init} end  
  P  
end
```
- Out is the final state when the agent terminates
  - It never terminates here, but in another definition it might

38

## Example Cell agent



- This agent behaves like a cell!

```
fun {CellProcess S M}  
  case M  
  of assign(New) then New  
  [] access(Old) then Old=S S  
  end  
end
```

- Cells and ports are equivalent in expressiveness
  - Even though they look very different

39

## Uniform interfaces (1)



- We can create and use a cell agent:

```
declare Cell  
Cell={NewPortObject CellProcess 0}  
{Send Cell assign(1)}  
local X in {Send Cell access(X)} {Browse X} end
```

- We want to have the same interface as objects:

```
{Cell assign(1)}  
local X in {Cell access(X)} {Browse X} end
```

40

## Uniform interfaces (2)



- We change the output to be a procedure:

```
fun {NewPortObject Init F}  
  P Out  
in  
  thread S in P={NewPort S} Out={FoldL S F Init} end  
  proc {$ M} {Send P M} end  
end
```

- P is hidden inside the procedure by lexical scoping
- This makes it easier to use port objects or standard objects as we saw before

41

## Message protocols



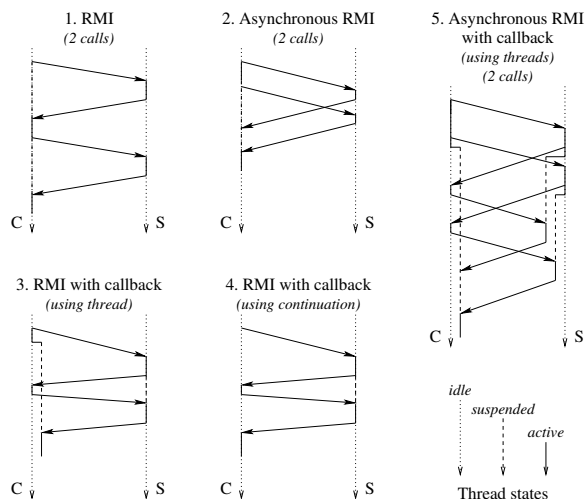
42

## Message protocols (1)

- A message protocols is a sequence of messages between two or more parties that can be understood at a higher level of abstraction than individual messages
- Using port objects, let us investigate some important message protocols
- Code given during the lecture!
  - Also in Section 5.3 of the course book

43

## Message protocols (2)



- We start with a simple RMI
- We then make it asynchronous and add callbacks
- The most complicated protocol is asynchronous RMI with callback


44

# Active objects



45

## Active objects (Section 7.8)



- An active object is a port object whose behavior is defined by a class
- Active objects combine the abilities of object-oriented programming (including polymorphism and inheritance) and message-passing concurrency
- To explain active objects, we refresh your memory on object-oriented programming and we introduce classes in Oz

46

## Classes and objects in Oz



- We saw objects in the course
- We now complete this explanation by introducing classes and their Oz syntax

```
class Counter
  attr i
  meth init(X)
    i := X
  end
  meth inc(X)
    i := @i + X
  end
  meth get(X)
    X=@i
  end
end
```

- Create an object:

```
Ctr={New Counter init(0)}
```

- Call the object:

```
{Ctr inc(10)}
{Ctr inc(5)}
local X in
  {Ctr get(X)}
  {Browse X}
end
```

47

## Defining active objects



- Active objects are defined by combining classes and port objects
- We use the uniform interface to make them look like standard Oz objects

```
fun {NewActive Class Init}
  Obj={New Class Init}
  P
in
  thread S in
    {NewPort S P}
    for M in S do {Obj M} end
  end
  proc {$ M} {Send P M} end
end
```

48




## Passive objects and active objects




- We make an important distinction between passive objects and active objects
- Standard objects in Oz (and many other languages, such as Java and Python) are now called **passive objects**
  - This is because they execute in the thread of their caller; they do not have their own thread
- This is in contrast to **active objects**, which have their own thread
- Let us compare passive and active objects!

49

## Concurrency comparison



- Passive objects cannot be safely called from more than one thread
- The method executions can overlap, which leads to concurrency bugs



- Active objects are completely safe when called from more than one thread
- The method executions are executed sequentially in the active object's own thread

50

## Passive objects are not concurrency-safe!



- The following code is buggy:

```
Ctr={New Counter init(0)}  
thread {Ctr inc(1)} end  
thread {Ctr inc(1)} end  
local X in  
  {Ctr get(X)}  
  {Browse X}  
end
```

- This can display 1! Why?
  - Look at the instruction `i := @i + 1`
  - If the scheduler puts T1 to sleep after `@i` and before `i :=`, executes T2 fully, and then resumes T1

- The following code is correct:

```
Ctr={NewActive Counter init(0)}  
thread {Ctr inc(1)} end  
thread {Ctr inc(1)} end  
local X in  
  {Ctr get(X)}  
  {Browse X}  
end
```

- This will always display 2
  - Because the two methods are executed sequentially by Ctr's thread

51

(to be continued)



52